

# CS182 project - Deliver ideas of CGCNN

## General Explanation

The key idea of CGCNN is to represent crystal structure by a crystal graph, which encodes not only the atomic information but also bonding information between atoms based on the distance. Here we will let each atomic feature vector as  $v_i$ , which encodes the property of the atom corresponding to node  $i$ . The edge feature vector is represented as  $u(i,j)_k$ , which encodes  $k$ -th bond connecting atom  $i, j$ . Note that there can be several bonds between atoms, which originate from the periodic nature of crystals.

Obviously, there are many ways to encode atomic and bond data. For the sake of simplicity, we will follow the methods used in the CGCNN paper, which used a pre-defined encoding vector to change atoms into vectors. These vectors have 92 dimensions, and for 100 atoms in the periodic table, they have different encodings which consist of 0 and 1. For the edge feature vectors, there also can be many ways to encode bond information; bond length, angle, and covalency. However, we will only use bond-length information between nearest neighbors, by applying a Gaussian kernel to change it to encoding vectors. One can follow the homework notebook on how this embedding is done to NaCl(salt) Crystal.

## Encoding Crystal into Crystal Graph

```
In [3]: # Note: Need pymatgen
```

```
In [2]: # Uncomment this to install pymatgen
        #!pip install pymatgen
```

```
In [ ]: import os
import sys
import csv
import json
import torch
import torch.nn as nn
import random
import warnings
warnings.filterwarnings('ignore')
import functools
import sklearn
import numpy as np
import networkx as nx
import matplotlib.pyplot as plt
from matplotlib.colors import to_rgba_array

from pymatgen.core.structure import Structure
from pymatgen.analysis.graphs import StructureGraph
from pymatgen.analysis.local_env import CrystalNN
from torch.utils.data import Dataset, DataLoader
from torch.utils.data.dataloader import default_collate
from torch.utils.data.sampler import SubsetRandomSampler
import torch.optim as optim
from torch.optim.lr_scheduler import MultiStepLR
```

```

from data_utils import CIFData
from data_utils import AtomCustomJSONInitializer
from data_utils import AtomInitializer
from data_utils import GaussianDistance

```

```

In [4]: # Uncomment this to extract data
        # !unzip cgcnn_data.zip
        # !unzip hw_data.zip

```

Let's convert salt (NaCl) to crystadddl graph.

You can use print method to see the lattice and position of Na, Cl atoms in the cell.

We will also visualize it to a graph. Each node of graph is atom in the crystal, and edge is bonds that have distance as information.

Note that we will use pymatgen package (Python Materials Genomics) which is an open source Python library for materials analysis. But you don't need to worry about that too much. Only important class is 'Structure' class, which can load structure data of materials by Structure.from\_file method. After you load it, you can get the site information using Structure\_object[index\_want\_to\_check], and this Site object has some useful attributes, such as specie. You can check what specie is in the site by using Structure\_object[index].specie

```

In [5]: # Load structure from CIF file
        nacl = Structure.from_file('hw_data/sample_regression/1000041.cif')
        print(nacl)

```

```

Full Formula (Na4 Cl4)
Reduced Formula: NaCl
abc      :   5.620000   5.620000   5.620000
angles:   90.000000   90.000000   90.000000
pbc      :          True          True          True
Sites (8)
  #  SP      a      b      c
---  ---  ---  ---  ---
  0  Na+    0      0      0
  1  Na+    0     0.5     0.5
  2  Na+    0.5    0     0.5
  3  Na+    0.5    0.5    0
  4  Cl-    0.5    0.5    0.5
  5  Cl-    0.5    0      0
  6  Cl-    0     0.5    0
  7  Cl-    0      0     0.5

```

```

In [6]: # Visualization of the structure using the CIF file
        crystal = CrystalNN()

        sg = StructureGraph.with_local_env_strategy(nacl, crystal)

        node_labels = {}
        for i, site in enumerate(nacl):
            label = "{}".format(site.specie)
            node_labels[i] = label

        edge_weights = {}
        for edge in sg.graph.edges():
            dist = nacl.get_distance(edge[0], edge[1])
            edge_weights[edge] = dist


        pos = nx.spring_layout(sg.graph)

```

```

nx.draw(sg.graph, pos, with_labels=True, labels=node_labels, font_size=12, node_
nx.draw_networkx_edge_labels(sg.graph, pos, edge_labels=edge_weights, font_size=
plt.axis("off")
plt.show()

```



Next we will encode atoms in the node to atomistic feature vectors using pre-defined atom embedding.

atom\_init.json is containing vector embedding of atoms, where key (1, 2, 3, ..., 100) represent atomic number and values are encoding vectors.

You can try different atom embedding too if you want.

```

In [ ]: # Load embedding file.
element_embedding_file = 'hw_data/sample_regression/atom_init.json'
with open(element_embedding_file) as f:
    elem_embedding = json.load(f)
elem_embedding = {int(key): value for key, value
                  in elem_embedding.items()}

atom_fea = None
#####
# TODO: Encode crsytal data to atomistic features.
# (Hint: Use np.vstack and elem_embedding to join all embedded specie numbers of
# Also, check out pymatgen.core.structure.specie.number attribute.
# nacl[i].specie will give you what kind of specie it is,
# And number attribute will give atomic number of that specie.)
# Atom_feature shd have shape of (# of atoms, len(embedding vector))

# YOUR CODE HERE

```

```
# Solution
atom_fea = np.vstack([elem_embedding[nacl[i].specie.number]
                      for i in range(len(nacl))])
#####
assert atom_fea is not None
atom_fea = torch.Tensor(atom_fea)

assert atom_fea.shape == (8, 92)
```

Next, we will get neighbor information from each atoms in the cell. We will get help from pymatgen package. get\_all\_neighbors function of structure object returns atoms within the input radius. Note that here len(all\_nbrs) is 8 since there are 8 atoms (4 Na<sup>+</sup>, 4 Cl<sup>-</sup>) in the cell. Each list contain the neighbor atom information considering periodicity. We will use 12 nearest neighbors after sorting with distance.

```
In [8]: all_nbrs = nacl.get_all_neighbors(r = 8, include_index=True)
all_nbrs = [sorted(nbrs, key=lambda x: x[1]) for nbrs in all_nbrs]

assert len(all_nbrs) == 8

nbr_fea_idx, nbr_fea = [], []
for nbr in all_nbrs:
    # Note: x[1] returns distance to neighbors.
    # Note: x[2] returns index of original structure object
    nbr_fea_idx.append(list(map(lambda x: x[2],
                                nbr[:12])))
    nbr_fea.append(list(map(lambda x: x[1],
                            nbr[:12])))

# nbr_fea_idx contain information of nearest neighbor atoms
# from ith row (ith atom in the cell)
# For example, 0th atom (Na+ (0.0000, 0.0000, 0.0000)) is neighbored
# with 5th, 6th, 7th, etc...
# nbr_fea contain information of nearest neighbor distance.
nbr_fea_idx, nbr_fea = np.array(nbr_fea_idx), np.array(nbr_fea)
nbr_fea_idx = torch.LongTensor(nbr_fea_idx)

nbr_fea_idx_standard = torch.LongTensor([[5, 6, 7, 7, 6, 5, 2, 1, 2, 1, 3, 3],
                                         [4, 7, 6, 7, 6, 4, 2, 0, 3, 2, 3, 3],
                                         [4, 7, 5, 7, 5, 4, 0, 3, 1, 3, 3, 1],
                                         [4, 6, 5, 6, 5, 4, 2, 1, 1, 2, 0, 0],
                                         [3, 1, 2, 3, 2, 1, 6, 5, 6, 5, 7, 7],
                                         [2, 3, 0, 3, 0, 2, 4, 4, 6, 6, 4, 7],
                                         [1, 3, 0, 3, 0, 1, 4, 4, 5, 5, 5, 4],
                                         [2, 1, 0, 2, 1, 0, 4, 5, 5, 4, 6, 6]])

assert torch.equal(nbr_fea_idx, nbr_fea_idx_standard)
```

Now we have two features, atomic feature and neighbor feature. Note that neighbor feature is discontinuous information with respect to the distance. Therefore we will expand (which means we will stack data based on the current discontinuous distance value) neighbor feature using Gaussian Kernel (or Gaussian filter) [https://en.wikipedia.org/wiki/Gaussian\\_filter](https://en.wikipedia.org/wiki/Gaussian_filter) While expanding the neighbor feature, one can set the distance range and step. Here we will use minimum distance 0, maximum distance 12, and step size of 0.2, which will expanded to dimension size of 61.

```
In [9]: dmin = 0
dmax = 12
step = 0.2
var = step
filter_step = np.arange(dmin, dmax+step, step)

def expand(distances):
    #####
    # TODO: Implement the expand function using a Gaussian kernel.
    # raise NotImplementedError()

    # Solution
    return np.exp(-(distances[..., np.newaxis] - filter_step)**2 / var**2)
    #####

nbr_fea_gaussian = expand(nbr_fea)
```

```
In [10]: gdf = GaussianDistance(dmin=0, dmax=12, step=0.2)
nbr_fea = gdf.expand(nbr_fea)
```

```
In [11]: assert np.array_equal(nbr_fea_gaussian, nbr_fea)
```

## Build a Model

After building the crystal graph, the convolutional neural network act on top of the graph, and consecutive fully connected networks and pooling layer are used to predict the property of crystals. The convolution layer iteratively updates the atom feature vector  $v_i$  by message passing with surrounding atoms and bonds with a non-linear graph convolution function:

$$v_i^{(t+1)} = \text{Conv}(v_i^{(t)}, v_j^{(t)}, u_{(i,j)_k}), (i, j)_k \in G$$

And then consecutive pooling layer and hidden fully connected layer will generate output of model:

$$v_c = \text{Pool}(v_0^{(0)}, v_1^{(0)}, \dots, v_N^{(0)}, \dots, v_N^{(R)})$$

Note that we can use any Convolution and Pooling layer if we maintain the permutation invariance. Here we will try two different convolution layer. Try to implement WeightShareConvLayer in the TODO section of model\_hw.py. We will offer solution for the ConvLayer. Note that we will do not change pooling layer, and you can use any activation function you use. But try to use sigmoid function here. After you implement WeightShareConvLayer, try to implement CrystalGraphConvNet by filling in the TODO section of model\_hw.py. You can find a detailed explanation there.

1. WeightShareConvLayer; This is a Convolution weight matrix is shared by all neighbors.

$$v_i^{(t+1)} = g[(\sum_{j,k} v_j^{(t)} \oplus u_{(i,j)_k}) W_c^{(t)} + v_i^{(t)} W_s^{(t)} + b^{(t)}]$$

Here  $W_c$ ,  $W_s$ , and  $b$  are the convolution weight matrix, self-weight matrix, and bias.

1. ConvLayer; This is a Convolution weight matrix is differentiated at all neighbor pairs.

$$v_i^{(t+1)} = v_i^{(t+1)} + \sum_{j,k} \sigma(z_{(i,j)_k}^{(t)} W_f^{(t)} + b_f^{(t)}) \odot g(z_{(i,j)_k}^{(t)} W_s^{(t)} + b_s^{(t)})$$

Here  $z_{(i,j)_k} = v_i \oplus v_j \oplus u_{(i,j)_k}$ ,  $\odot$  denotes element-wise multiplication and  $\sigma$  denotes a sigmoid function.

```
In [12]: from model_hw import CrystalGraphConvNet
         from model_hw import ConvLayer
         from model_hw import WeightShareConvLayer
```

```
In [13]: # For testing purpose
         data_dir = './cgcn_data/sample-regression'
         test = CIFData(data_dir)
         (atom_fea, nbr_fea, nbr_fea_idx), target, cif_id = test[-1]
         structures, _, _ = test[-1]
         orig_atom_fea_len = structures[0].shape[-1]
         nbr_fea_len = structures[1].shape[-1]
```

```
In [14]: torch.manual_seed(123)

         convlayer_test = ConvLayer(92, 41)
         convlayer_test_result = convlayer_test.forward(atom_fea, nbr_fea, nbr_fea_idx)

         #####
         # TODO: Implement the WeightShareConvLayer model in model_hw.py
         #####

         weightshare_test = WeightShareConvLayer(92, 41)
         weightshare_test_result = weightshare_test.forward(atom_fea, nbr_fea, nbr_fea_idx)

         assert convlayer_test_result.shape == (8, 92)
         assert weightshare_test_result.shape == (8, 92)
```

```
In [15]: # Simple test
         # Set seed using manual_seed
         # TODO explanation about crystal_atom_idx
         torch.manual_seed(123)
         crystal_atom_idx = [torch.Tensor([0, 1, 2, 3, 4, 5, 6, 7]).long()]
         #####
         # TODO: Implement the CrystalGraphConvNet model in model_hw.py
         #####

         model1 = CrystalGraphConvNet(orig_atom_fea_len, nbr_fea_len, option='C')
         out = model1.forward(atom_fea, nbr_fea, nbr_fea_idx, crystal_atom_idx)
         assert (out.item() - 0.6627038) < 1e-6

         model2 = CrystalGraphConvNet(orig_atom_fea_len, nbr_fea_len, option='WC')
         out = model2.forward(atom_fea, nbr_fea, nbr_fea_idx, crystal_atom_idx)
         assert (out.item() - -0.5442343) < 1e-6
```

## Training

Now we are all set to test our model. We will train our model to simplified dataset, which consists of ~700 Perovskite materials. We will predict the energy (which is calculated already with DFT, unit is eV) of Perovskites, and see how the model performs. After run the following

section, try to answer the question. Note that you can find your prediction and corresponding structures in test\_result.csv file.

Q. If you built a model, let's train and visualize it. What is the difference between two models? Which one performs better and why do you think so?

Based on what you observed, can you make a convolution layer that performs better?

Hint: Among two convolution layers, ConvLayer performs better in general. This is because the local environments can always change, so maintaining same weights to all local environments makes prediction worse. However, you can see that in Perovskite dataset, WeightShareConvLayer performs better because Perovskite all have similar symmetry, and local environments of them are not much different

```
In [16]: from data_utils import collate_pool, get_train_val_test_loader
from train_utils import Normalizer, train, validate, save_checkpoint
from model_hw import CrystalGraphConvNet
from random import sample
```

```
In [17]: torch.manual_seed(123)

# set parameters
data_dir = './hw_data/perovskite_energy'
batch_size = 8
train_ratio, val_ratio, test_ratio = 0.8, 0.1, 0.1

# get dataset

dataset = CIFData(data_dir)
collate_fn = collate_pool
train_loader, val_loader, test_loader = get_train_val_test_loader(
    dataset=dataset,
    collate_fn=collate_fn,
    batch_size=batch_size,
    train_ratio=train_ratio,
    val_ratio=val_ratio,
    test_ratio=test_ratio,
    return_test=True)
```

```
In [18]: # normalize target

if len(dataset) < 500:
    warnings.warn('Dataset has less than 500 data points. '
                  'Lower accuracy is expected. ')
    sample_data_list = [dataset[i] for i in range(len(dataset))]
else:
    sample_data_list = [dataset[i] for i in
                        sample(range(len(dataset)), 500)]
_, sample_target, _ = collate_pool(sample_data_list)
normalizer = Normalizer(sample_target)
```

```
In [19]: # build model
structures, _, _ = dataset[0]
orig_atom_fea_len = structures[0].shape[-1]
nbr_fea_len = structures[1].shape[-1]
#####
```

```

# TODO: Tune the following parameters and model to optimize performance.
# Note: You can change your convolution layer with option argument in
# CrystalGraphConvNet. C is ConvLayer, WC is WeightShareConvLayer.
#####
# number of hidden atom features in conv layers
atom_fea_len = 64
# number of hidden features after pooling
h_fea_len = 128
# number of conv layers
n_conv = 3
# number of hidden layers after pooling
n_h = 1

model = CrystalGraphConvNet(orig_atom_fea_len, nbr_fea_len,
                             atom_fea_len=atom_fea_len,
                             n_conv=n_conv,
                             h_fea_len=h_fea_len,
                             n_h=n_h, option='C')

```

```

In [20]: # set hyperparameters, change the parameters to get better results
epochs = 15
criterion = nn.MSELoss()
lr = 0.01
momentum = 0.9
weight_decay = 0

optimizer = optim.SGD(model.parameters(), lr,
                       momentum=momentum,
                       weight_decay=weight_decay)

# optimizer = optim.Adam(model.parameters(), lr,
#                          weight_decay=weight_decay)
lr_milestones = [100]
scheduler = MultiStepLR(optimizer, milestones=lr_milestones,
                         gamma=0.1)

```

```

In [21]: running_record = {'train_loss': [], 'train_mae': [],
                           'val_loss': [], 'val_mae': []}

for epoch in range(epochs):
    print('Epoch {}/{}'.format(epoch, epochs - 1))
    print('-' * 10)

    best_mae_error = 1e10
    # train for one epoch
    train_loss, train_mae = train(train_loader, model, criterion, optimizer, epo

    # evaluate on validation set
    val_loss, val_mae = validate(val_loader, model, criterion, normalizer)

    # append loss and mae to running record. Convert tensor to float if necessar
    running_record['train_loss'].append(train_loss)
    running_record['train_mae'].append(train_mae)
    running_record['val_loss'].append(val_loss)
    running_record['val_mae'].append(val_mae)

    if val_mae != val_mae:
        print('Exit due to NaN')
        sys.exit(1)

```



```

scheduler.step()

# remember the best mae_eror and save checkpoint
is_best = val_mae < best_mae_error
best_mae_error = min(val_mae, best_mae_error)

save_checkpoint({
    'epoch': epoch + 1,
    'state_dict': model.state_dict(),
    'best_mae_error': best_mae_error,
    'optimizer': optimizer.state_dict(),
    'normalizer': normalizer.state_dict(),
}, is_best)

```

Epoch 0/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:21<00:00, 3.12it/s]
Train: Time 0.246      Data 0.168      Loss 1.0064      MAE 54.459
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:08<00:00, 1.03it/s]
Test:  Time 0.412      Loss 0.4892      MAE 40.253

```

Epoch 1/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:19<00:00, 3.40it/s]
Train: Time 0.220      Data 0.152      Loss 0.6307      MAE 39.189
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:08<00:00, 1.04it/s]
Test:  Time 0.405      Loss 0.8308      MAE 51.584

```

Epoch 2/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:18<00:00, 3.54it/s]
Train: Time 0.208      Data 0.150      Loss 0.6201      MAE 38.471
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:08<00:00, 1.04it/s]
Test:  Time 0.406      Loss 0.4942      MAE 35.135

```

Epoch 3/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:19<00:00, 3.49it/s]
Train: Time 0.212      Data 0.149      Loss 0.5625      MAE 36.203
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:13<00:00, 1.46s/it]
Test:  Time 0.904      Loss 0.5425      MAE 40.618

```

Epoch 4/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:19<00:00, 3.50it/s]
Train: Time 0.211      Data 0.150      Loss 0.6808      MAE 43.837
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:08<00:00, 1.06it/s]
Test:  Time 0.389      Loss 0.4754      MAE 41.545

```

Epoch 5/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:19<00:00, 3.36it/s]
Train: Time 0.223      Data 0.156      Loss 0.5892      MAE 37.267
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:08<00:00, 1.05it/s]
Test:  Time 0.395      Loss 0.4730      MAE 44.930

```

Epoch 6/14

-----

```

100%|████████████████████████████████████████████████████████████████████████████████| 67/67 [00:19<00:00, 3.46it/s]
Train: Time 0.214      Data 0.153      Loss 0.4989      MAE 34.168
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:09<00:00, 1.01s/it]
Test:  Time 0.458      Loss 0.6255      MAE 43.454

```

Epoch 7/14

-----

```

100%|████████████████████████████████████████| 67/67 [00:19<00:00, 3.51it/s]
Train: Time 0.210 Data 0.151 Loss 0.5025 MAE 33.486
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.05it/s]
Test: Time 0.405 Loss 0.4627 MAE 42.372
Epoch 8/14
-----

100%|████████████████████████████████████████| 67/67 [00:18<00:00, 3.56it/s]
Train: Time 0.206 Data 0.149 Loss 0.5437 MAE 37.471
100%|████████████████████████████████████████| 9/9 [00:09<00:00, 1.04s/it]
Test: Time 0.484 Loss 0.3625 MAE 34.471
Epoch 9/14
-----

100%|████████████████████████████████████████| 67/67 [00:18<00:00, 3.55it/s]
Train: Time 0.207 Data 0.149 Loss 0.4624 MAE 34.295
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.06it/s]
Test: Time 0.393 Loss 0.5161 MAE 44.592
Epoch 10/14
-----

100%|████████████████████████████████████████| 67/67 [00:18<00:00, 3.62it/s]
Train: Time 0.202 Data 0.142 Loss 0.4524 MAE 34.097
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.05it/s]
Test: Time 0.401 Loss 0.4711 MAE 38.106
Epoch 11/14
-----

100%|████████████████████████████████████████| 67/67 [00:20<00:00, 3.35it/s]
Train: Time 0.224 Data 0.159 Loss 0.4972 MAE 34.538
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.05it/s]
Test: Time 0.401 Loss 0.4282 MAE 41.408
Epoch 12/14
-----

100%|████████████████████████████████████████| 67/67 [00:18<00:00, 3.59it/s]
Train: Time 0.204 Data 0.147 Loss 0.4259 MAE 33.707
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.06it/s]
Test: Time 0.392 Loss 0.4729 MAE 39.691
Epoch 13/14
-----

100%|████████████████████████████████████████| 67/67 [00:19<00:00, 3.52it/s]
Train: Time 0.209 Data 0.149 Loss 0.4054 MAE 32.458
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.05it/s]
Test: Time 0.399 Loss 0.4873 MAE 34.707
Epoch 14/14
-----

100%|████████████████████████████████████████| 67/67 [00:18<00:00, 3.64it/s]
Train: Time 0.200 Data 0.138 Loss 0.4295 MAE 32.980
100%|████████████████████████████████████████| 9/9 [00:08<00:00, 1.05it/s]
Test: Time 0.398 Loss 0.5460 MAE 44.589

```

```

In [22]: # test best model
print('-----Evaluate Model on Test Set-----')
best_checkpoint = torch.load('model_best.pth.tar')
model.load_state_dict(best_checkpoint['state_dict'])
_, l=validate(test_loader, model, criterion, normalizer, test=True)

```

```

-----Evaluate Model on Test Set-----
100%|████████████████████████████████████████| 9/9 [00:09<00:00, 1.01s/it]
Test: Time 0.459 Loss 1.2363 MAE 53.165

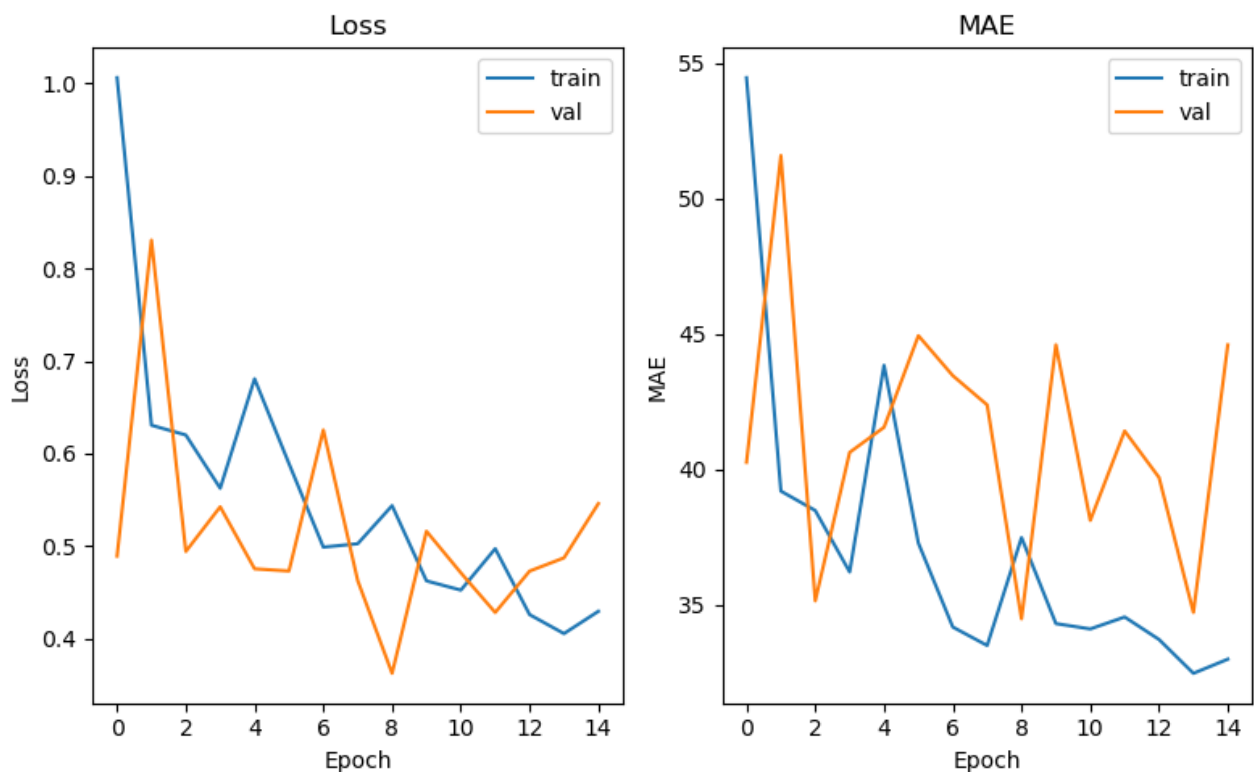
```

```
In [23]: # visualize the training and val loss
# visualize the training and val mae
import matplotlib.pyplot as plt
from matplotlib.ticker import MaxNLocator

fig, ax = plt.subplots(1, 2, figsize=(8, 5))
ax[0].plot(running_record['train_loss'], label='train')
ax[0].plot(running_record['val_loss'], label='val')
ax[0].set_xlabel('Epoch')
ax[0].set_ylabel('Loss')
ax[0].legend()
ax[0].set_title('Loss')
# set x axis to be integer
ax[0].xaxis.set_major_locator(MaxNLocator(integer=True))

ax[1].plot(running_record['train_mae'], label='train')
ax[1].plot(running_record['val_mae'], label='val')
ax[1].set_xlabel('Epoch')
ax[1].set_ylabel('MAE')
ax[1].legend()
ax[1].set_title('MAE')
# set x axis to be integer
ax[1].xaxis.set_major_locator(MaxNLocator(integer=True))

plt.tight_layout()
plt.show()
```



In [ ]: